

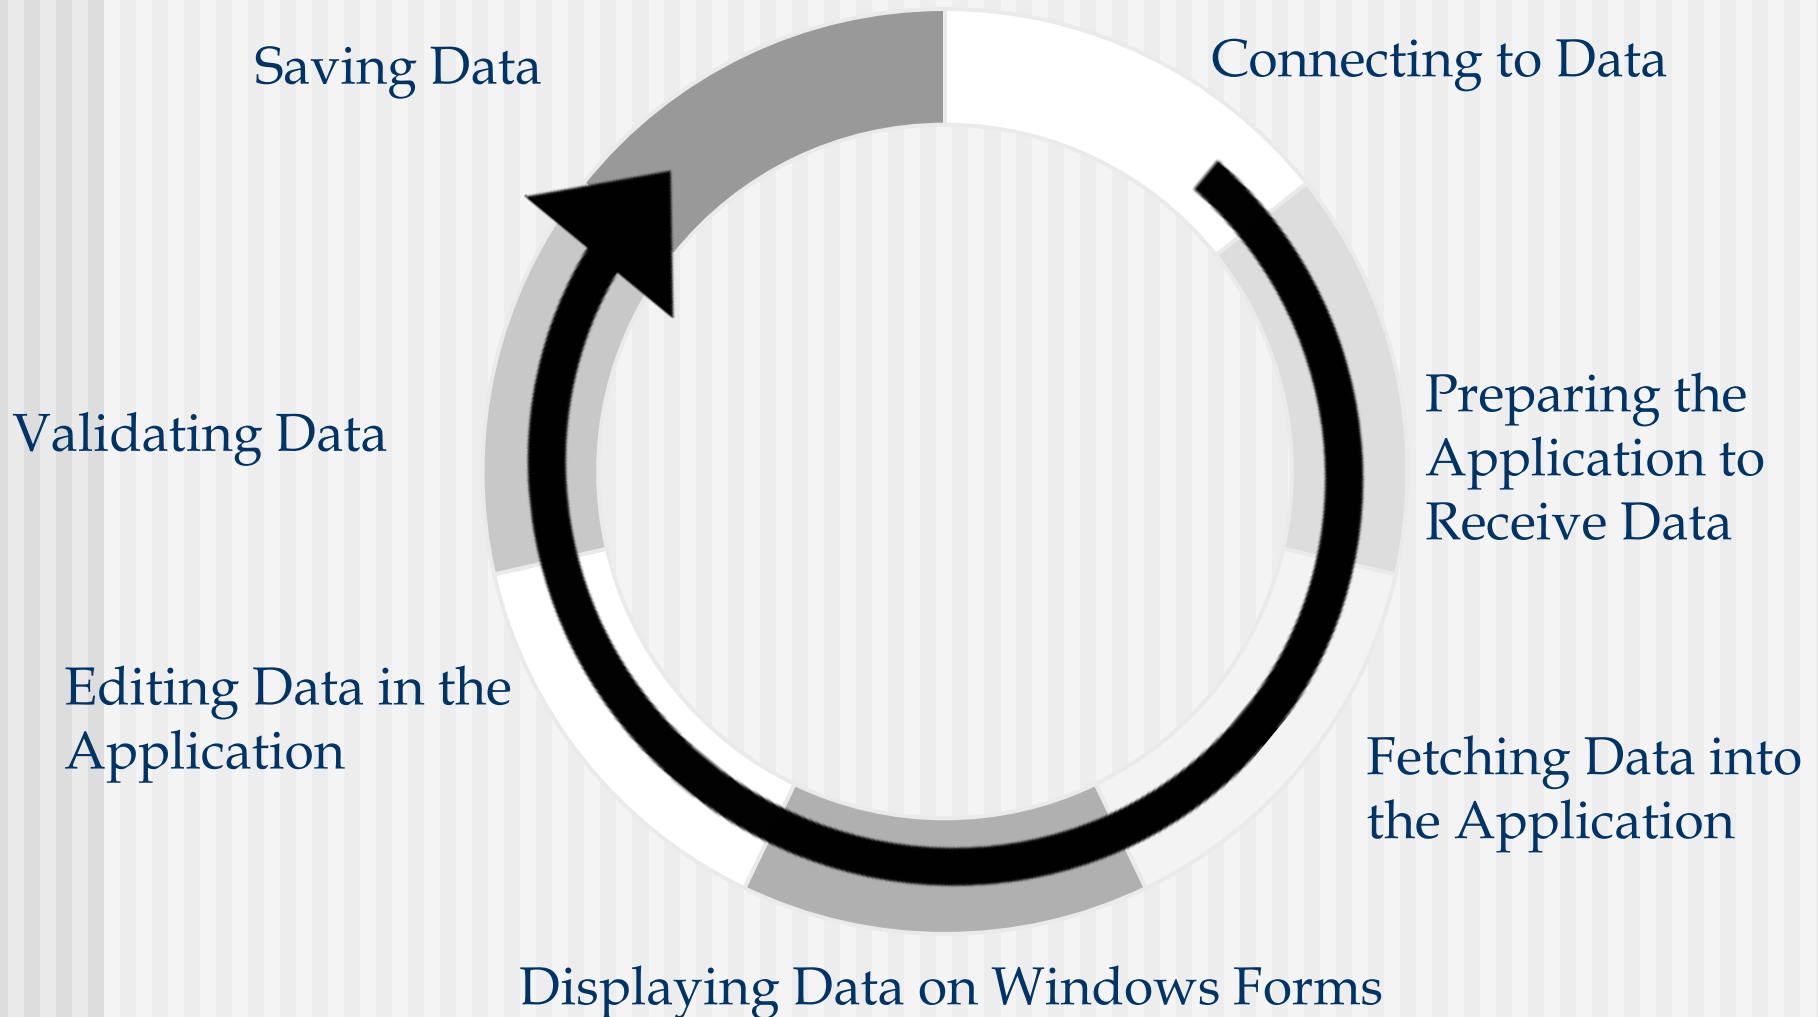
# Seminar 1

---

ADO.NET

# The Data Cycle

---



# The Data Cycle

---

- connecting to data: establish a two-way communication between the application and the database server (e.g., *TableAdapter*)
- preparing the app to receive data: when using a disconnected data model there are specific objects that temporarily store data (e.g., *datasets*, *LINQ to SQL* objects)
- fetching data: execute queries and stored procedures (e.g., by using *TableAdapters*)

# The Data Cycle

---

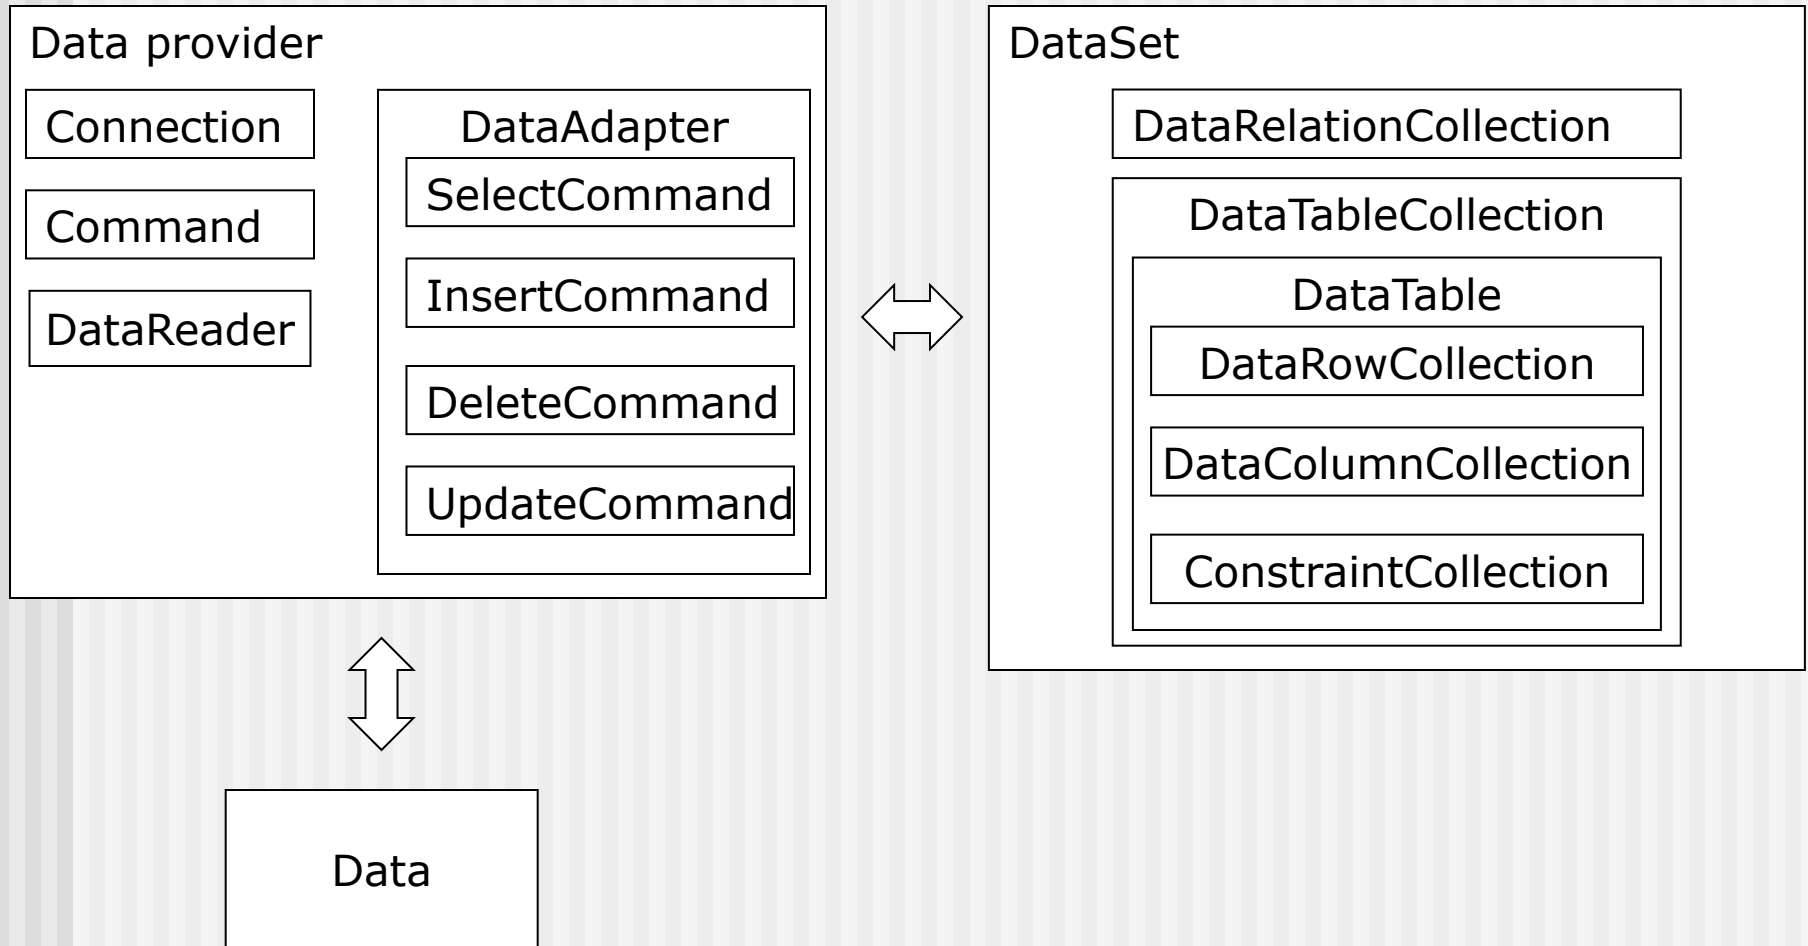
- displaying data: data-bound controls
- editing and validating data: add / modify / delete records; verify if the new values meet the requirements of the application
- saving data: send changes back to the database (e.g., *TableAdapterManager*)

# Data Models

---

- typed / untyped datasets
- conceptual model based on the *Entity Data Model*
- *LINQ to SQL*

# ADO.NET



# Datasets

---

- System

- Data

- DataSet
    - DataTable
    - DataRelation
    - DataColumn
    - DataRow
    - Constraint

# Datasets - the DataSet class

---

- object containing data tables that can temporarily store the data used in the application
- typed / untyped
- local in-memory cache
- also works when the application disconnects from the database
- structure similar to that of a relational database (tables, rows, columns, constraints, relationships)



# Datasets - the DataSet class

---

- properties:

- *Tables* (*DataTableCollection*) – collection of tables in the DataSet
- *Relations* (*DataRelationCollection*) – collection of relations (child / parent tables)

- methods:

- *Clear()* – clears all rows in all tables
- *HasChanges()* – indicates whether there are new / deleted / modified rows

# Datasets - the DataTable class

---

- properties:
  - *Rows (DataRowCollection)*
  - *Columns (DataColumnCollection)*
  - *ChildRelations and ParentRelations (DataRelationCollection)*

# Datasets - the DataRow class

---

- properties:
  - *RowState* - possible values: *Added*, *Deleted*, *Modified*, *Unchanged*

# SqlConnection

---

- represents a connection to a SQL Server database
- cannot be inherited
- if the SqlConnection goes out of scope, it is not closed => must ensure the connection is always closed (e.g., *Close*)
- properties:
  - *ConnectionString* - string used to open a SQL Server database  
(<http://www.connectionstrings.com/> )
  - *ConnectionTimeout* - time to wait to establish a connection before terminating the attempt and generating an error

# SqlConnection

---

- methods:
  - *Open()*
  - *Close()*
- if a *SqlException* is generated, the *SqlConnection* remains open when the severity level  $\leq 19$

# SqlCommand

---

- represents a Transact-SQL statement or stored procedure to be executed on a SQL Server database
- cannot be inherited
- properties:
  - *CommandText*
  - *CommandTimeout*
- methods:
  - *ExecuteNonQuery* – returns the number of affected rows

# SqlCommand

---

- methods:
  - *ExecuteScalar* – returns the first column of the first row in the answer set
  - *ExecuteReader* – builds a SqlDataReader

## SqlDataReader

- reads a forward-only stream of rows from a SQL Server database

# SqlDataAdapter

---

- bridge between a DataSet and SQL Server to obtain data and save changes back to the database
- a set of commands and a database connection
- properties:
  - *UpdateCommand* - statement / stored procedure used to update records in the data source
  - *InsertCommand* - statement / stored procedure used when inserting records into the data source
  - *DeleteCommand* - statement / stored procedure used to delete records



# SqlDataAdapter

---

- methods:

- *Fill(DataSet, String)* - adds or refreshes rows in the DataSet object to match those in the data source (2<sup>nd</sup> param – name of table)
- *Update(DataSet, String)* – changes the values in the database by executing INSERT / UPDATE / DELETE statements for every added / modified / deleted row

# Console

---

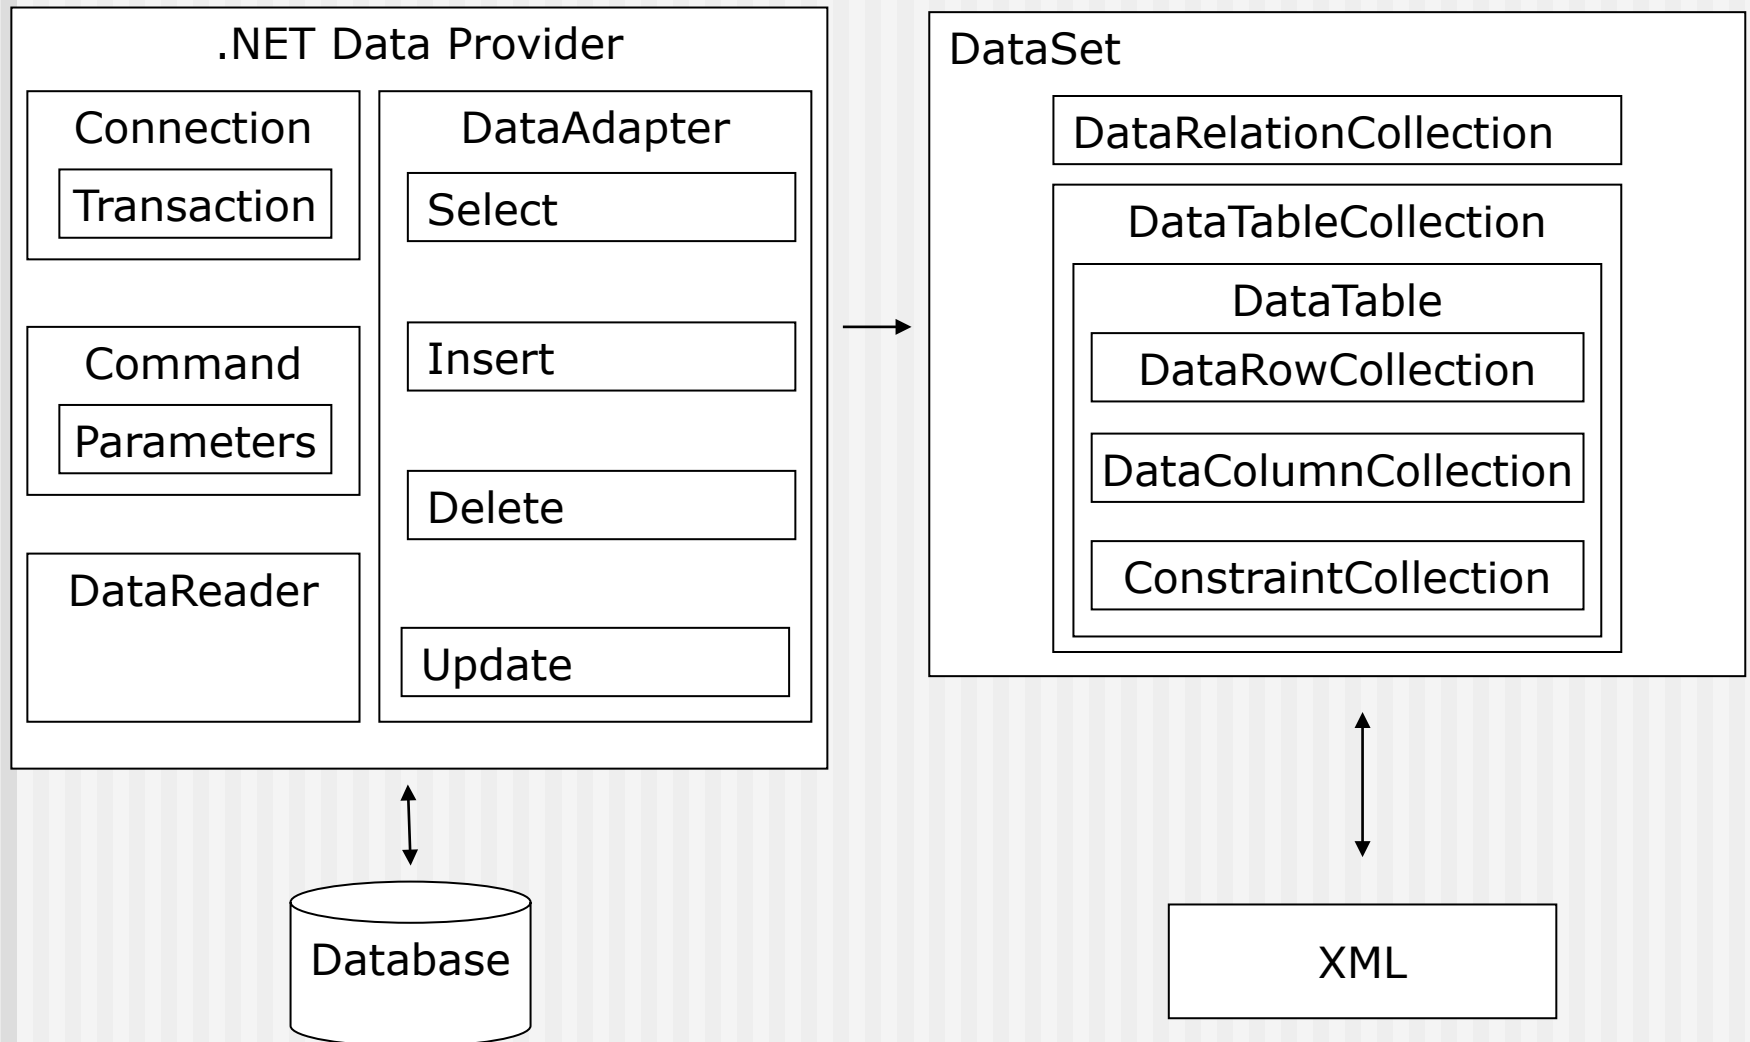
- it represents the standard input, output, error streams for console applications
- properties:
  - *WindowLeft, WindowTop, WindowHeight, WindowWidth, BackgroundColor, Title, etc.*
- methods:
  - *Write(...), WriteLine(...)*
  - *Read(), ReadLine(), ReadKey()*

# Seminar 2

---



## ADO.NET Data Binding

# The ADO.NET Architecture



# Data Binding – Windows Forms

Windows Form



# Windows Forms – Structures to Bind To

---

- list-based binding – the object should support the *ICollection* interface
- ADO.NET provides data structures suitable for data binding
  - *DataColumn*
  - *DataTable* – with columns, rows, constraints
  - *DataRowView* – customized view of a single data table
  - *DataSet* – with tables, relationships
  - *DataRowManager* – customized view of a *DataSet*

# *CurrencyManager*

---

- keeps data-bound controls synchronized with each other
- for each data source associated with a Windows Form, there is one *CurrencyManager* object
- *currency* = the current position within a data structure
- the *Position* property - determines the current position of all the controls using the same *CurrencyManager* (the position within the underlying list)

# *BindingContext*

---

- manages the collection of *CurrencyManager* objects for any container control / form



# The *DataBindings* Property and the *Binding* Class

---

## ■ *DataBindings*

- property of a control
- retrieves the data bindings for the control
- one can bind any property of a control to the property of an object

## ■ *Binding* class

- the simple binding between the property value of an object and the property value of a control
- *Binding* (*String propertyName*, *Object dataSource*, *String dataMember*)

# Dataset – *Fill, Update*

- a dataset doesn't contain data by default
- data tables are filled with data when executing *TableAdapter* / data adapter (*SqlDataAdapter*) queries / commands

```
aTableAdapter.Fill(aDataSet.TableName);
```

- saving data:

```
aTableAdapter.Update(aDataSet.TableName);
```

- when the *Update* method is called, the value of the *RowState* property is examined to determine which records must be saved and which command (*InsertCommand*, *UpdateCommand*, *DeleteCommand*) must be executed

# Accessing Records

---

- every table exposes a collection of rows
- rows can be accessed through the collection's index or by using collection-specific statements in the host programming language

- typed dataset

```
TextBox1.Text = ds1.TableName[3].aField;
```

- untyped dataset

```
string val = (string)  
ds1.Tables["TableName"].Rows[0]["aField"];
```

# Related Tables, *DataRelation* Objects

- the data in a dataset's tables can be interrelated
- *DataRelation* objects can be created to describe the relationships among the dataset's tables
- a *DataRelation* object can be used to locate related records:
  - the *GetChildRows* method
    - called on a *DataRow* in the parent table
    - returns an array of related child records (*DataRow* objects)

# Related Tables, *DataRelation* Objects

- a *DataRelation* object can be used to locate related records:
  - the *GetParentRow* method
    - called on a *DataRow* in the child table
    - returns a single *DataRow* from the parent table

# Related Tables, *DataRelation* Objects

- return the child records for a parent record

```
string custID = "VINET";  
NorthwindDataSet.OrdersRow[] orders;  
orders = (NorthwindDataSet.OrdersRow[])  
    northwindDataSet.Customers.  
        FindByCustomerID(custID).GetChildRows  
            ("FK_Orders_Customers");  
MessageBox.Show(orders.Length.ToString());
```

# Related Tables, *DataRelation* Objects

- return the parent record for a child record

```
int orderID = 10835;  
NorthwindDataSet.CustomersRow customer;  
customer = (NorthwindDataSet.CustomersRow)  
    northwindDataSet.Orders.  
        FindByOrderID(orderID).  
        GetParentRow("FK_Orders_Customers");  
MessageBox.Show(customer.CompanyName);
```

# Creating an Application in Visual Studio

---

- create a new *Windows Forms* project
  - *File -> New -> Project*
  - select *Windows Forms App*, specify a name, choose a location
  - click on *OK*
  - the project is created and added to the *Solution Explorer*



# Creating an Application in Visual Studio

- create a *Data Source*
  - start the *Data Source Configuration Wizard* (*Data Sources* window)
  - choose a Data Source Type (e.g., Database)
  - choose a Database Model (Dataset)
  - choose the data connection
  - select the database objects (e.g., the required tables)

# Creating an Application in Visual Studio

- drag items (e.g., a particular table) from the Data Sources window onto the form to create *data-bound controls*

=> the following components are now visible in the component tray:

- *DataSet* - typed dataset that contains tables
- *BindingSource* - binds the controls on the form to the table in the dataset
- *BindingNavigator* - allows the user to navigate through the rows in the table
- *TableAdapter* - communication between the database and the dataset

# Creating an Application in Visual Studio

- drag items (e.g., a particular table) from the Data Sources window onto the form to create *data-bound controls*

=> the following components are now visible in the component tray:

- *TableAdapterManager* - controls the order of individual inserts, updates, and deletes

# Constraints

---

- two types of constraints: unique / foreign key
- unique constraint
  - all values in a set of columns must be unique
  - class *UniqueConstraint*
- foreign key constraint
  - defines rules on how to change related child records when a parent record is updated or deleted
  - class *ForeignKeyConstraint*

# Constraints

---

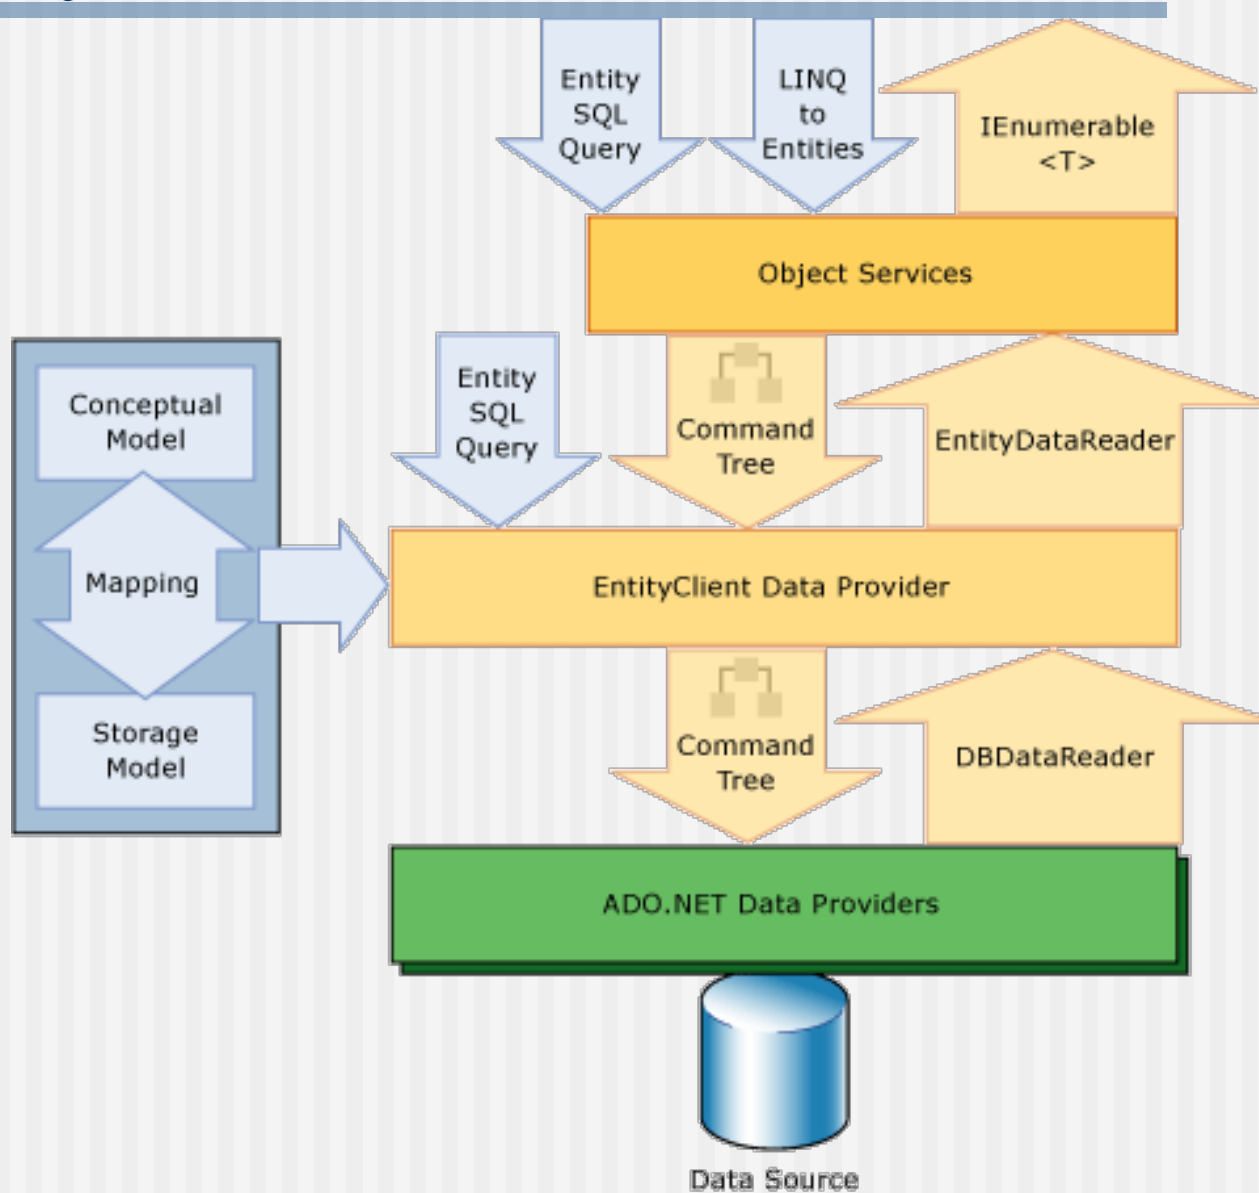
- a foreign key constraint is automatically added when creating a *DataRelation* object in a dataset
- a table's constraints can be retrieved using the *Constraints* property
- the boolean property *EnforceConstraints* in the *Dataset* class indicates whether constraints are enforced or not (by default it's *true*)

# Entity Framework

---

- set of technologies in ADO.NET that support the development of data-oriented applications
- developers can work with domain-specific objects and properties (as opposed to dealing with tables and columns)
- developers can query entities and relationships in the domain model; the Entity Framework translates such operations to data source-specific commands

# Entity Framework



# Seminar 3

---

## Transactions Concurrency Control in SQL Server



# Transactions in SQL Server

---

- combine multiple operations into a single unit of work
- the actions of each user are processed using a different transaction
- objective:
  - maximize throughput => transactions must be allowed to execute in parallel
- ACID properties
- serializability

# Transactions in SQL Server

- transaction invocation - mechanisms:
  - unless specified otherwise, each command is a transaction
  - BEGIN TRAN, ROLLBACK TRAN, COMMIT TRAN
  - SET IMPLICIT\_TRANSACTIONS ON
    - enables chained transactions
- SET XACT\_ABORT ON
  - SQL errors => rollback transaction

# Transactions in SQL Server

- local transactions / distributed transactions
- one can *nest* transactions (but transactions are not really nested)
- named *savepoints*
  - allow a portion of work in a transaction to be rolled back

# Concurrency Problems

- transaction isolation tackles four major concurrency problems:
  - *lost updates* - two transactions (writers) modify the same piece of data
  - *dirty reads* - a transaction (reader) reads uncommitted data, i.e., data changed by another ongoing transaction
  - *unrepeatable reads* - a row read by a transaction (reader) is changed by another transaction while the reader is in progress (if the 1<sup>st</sup> transaction reads the row again it will get different row values)

# Concurrency Problems

---

- transaction isolation tackles four major concurrency problems:
  - *phantoms* - transaction T1 reads a set of rows based on a search predicate; transaction T2 generates a new row (I/U) that matches the search predicate while T1 is ongoing; if T1 issues the same read operation, it will get an extra row

# Concurrency Problems

- transaction isolation is achieved through the locking mechanism
- *write locks*
  - exclusive locks, i.e., they don't allow other readers / writers
- *read locks*
  - allow other readers
  - don't allow other writers

# Concurrency Problems

---

- isolation levels determine:
  - whether read locks are acquired for read operations
  - the duration of read locks
  - whether key-range locks are acquired to prevent phantoms

# Locking in SQL Server

---

- locks
  - usually managed by the Lock Manager (not via apps)
- lock granularity:
  - *Row / Key, Page, Table, Extent\*, Database*
- hierarchy of related locks
  - locks can be acquired at several levels
- lock escalation
  - > 5000 locks per object (pros & cons)

\* contiguous group of 8 pages  
9



# Locking in SQL Server

- lock types:
- *Shared (S)*
  - read operations
- *Update (U)*
  - deadlock avoidance mechanism
- *Exclusive (X)*
  - write operations
  - incompatible with other locks

	S	X
S	Yes	No
X	No	No

# Locking in SQL Server

---

- lock types:
- *Exclusive (X)*
  - read operations by other transactions can be performed only when using the NOLOCK hint or the READ UNCOMMITTED isolation level
  - a transaction always acquires exclusive locks to modify data (regardless of the isolation level)
  - exclusive locks are released when the transaction completes execution

# Locking in SQL Server

---

- lock types:
- *Intent* (IX, IS, SIX)
  - intention to lock (for performance improvement purposes)
- *Schema* (Sch-M, Sch-S)
  - schema modification, schema stability
  - Sch-M
    - prevents concurrent access to the table
  - Sch-S
    - doesn't allow DDL operations to be performed on the table

# Locking in SQL Server

---

- lock types:
- *Bulk Update* (BU)
  - bulk load data concurrently into the same table
  - BULK INSERT statement
  - TABLOCK hint
- *Key-Range*
  - protect a range of rows implicitly included in a set of records read by a transaction (under the SERIALIZABLE isolation level)

# Key-Range Locking

---

- lock sets of rows defined by a predicate  
    ...**WHERE grade between 8 and 10**
- lock existing data, as well as data that doesn't exist
- use predicate “**grade between 8 and 10**” 2 times => obtain the same rows

# Transaction Workspace Locks

---

- every connection to a database acquires a *Shared\_Transaction\_Workspace* lock
- exceptions - connections to master, tempdb
- used to prevent:
  - DROP
  - RESTORE

# Isolation Levels in SQL Server

---

- **READ UNCOMMITTED**
  - allows dirty reads (a transaction can see uncommitted changes made by another ongoing transaction)
  - no S locks when reading data
- **READ COMMITTED** (default isolation level)
  - a transaction cannot read data that has been modified by another ongoing transaction
  - allows unrepeatable reads
  - S locks - released as soon as the SELECT operation is performed

# Isolation Levels in SQL Server

---

- **READ COMMITTED**
  - X locks - released at the end of the transaction
- **REPEATABLE READ**
  - holds S locks and X locks until the end of the transaction
  - doesn't allow dirty reads, unrepeatable reads
  - phantom reads can occur



# Isolation Levels in SQL Server

---

- **SERIALIZABLE**
  - highest isolation level
  - holds locks (including key-range locks) during the entire transaction
  - doesn't allow dirty reads, unrepeatable reads, phantom reads
- **SNAPSHOT**
  - working on a snapshot of the data
- SQL syntax
  - **SET TRANSACTION ISOLATION LEVEL ...**

# Isolation Levels in SQL Server

concurrency probl. / isolation level	Chaos	Read Uncommitted	Read Committed	Repeatable Read	Serializable
Lost Updates?	Yes	No	No	No	No
Dirty Reads?	Yes	Yes	No	No	No
Unrepeatable Reads?	Yes	Yes	Yes	No	No
Phantoms?	Yes	Yes	Yes	Yes	No

# Deadlocks

- SQL Server uses deadlock detection
- the transaction that's least expensive to roll back is terminated
- capture and handle error 1205
- SET LOCK\_TIMEOUT
  - specify how long (in milliseconds) a transaction waits for a locked resource to be released
  - value 0 - immediate termination
- SET DEADLOCK\_PRIORITY
  - values: *{LOW, NORMAL, HIGH, <numeric-priority>}*
  - *<numeric-priority> ::= {-10, -9, ..., 10}*

# Reduce the Likelihood of Deadlocks

---

- transactions - short & in a single batch
- obtain / verify input data from the user before opening a transaction
- access resources in the same order
- use a lower / a row versioning isolation level

# Seminar 4

---

## Multiversioning

# Monitoring Locks

---

- SQL Server Profiler
- sp\_lock
- sys.dm\_tran\_locks
- sys.dm\_tran\_active\_transactions

# Resource Types

---

- RID – row in a heap
- Key – row in an index
- Page
- HoBT – heap or B-tree
- Object – table, view, etc.
- File
- Database
- Metadata
- Application

# Isolation Levels in SQL Server

---

- **READ UNCOMMITTED**
  - allows dirty reads (a transaction can see uncommitted changes made by another ongoing transaction)
  - no S locks when reading data
- **READ COMMITTED** (default isolation level)
  - a transaction cannot read data that has been modified by another ongoing transaction
  - allows unrepeatable reads
  - S locks - released as soon as the SELECT operation is performed



# Isolation Levels in SQL Server

---

- **READ COMMITTED**
  - X locks - released at the end of the transaction
- **REPEATABLE READ**
  - holds S locks and X locks until the end of the transaction
  - doesn't allow dirty reads, unrepeatable reads
  - phantom reads can occur

# Isolation Levels in SQL Server

---

- **SERIALIZABLE**
  - highest isolation level
  - holds locks (including key-range locks) during the entire transaction
  - doesn't allow dirty reads, unrepeatable reads, phantom reads
- **SNAPSHOT**
  - working on a snapshot of the data
- SQL syntax
  - **SET TRANSACTION ISOLATION LEVEL ...**

# Multiversioning

---

- in a DBMS with multiversioning, every write operation on a data object  $O$  results in a new copy (i.e., *version*) of  $O$
- every time object  $O$  is read, the DBMS picks one of  $O$ 's versions
- writes do not overwrite each other

# Row-Level Versioning (RLV)

---

- introduced in SQL Server 2005
- useful when the user needs *committed* data (not necessarily *the most recent version* of the data)
- Read Committed Snapshot Isolation & Full Snapshot Isolation:
  - the reader never blocks; instead, it obtains data that has been previously committed
- the tempdb database stores the older versions of the data:
  - a snapshot of the database can be assembled using these old(er) versions

# Read Committed Snapshot Isolation

---

```
ALTER DATABASE MyDatabase  
SET READ_COMMITTED_SNAPSHOT ON
```

- operations see the most recent committed data as of the beginning of their execution =>
  - snapshot of the data at the command level
  - consistent reads at the command level
  - READ COMMITTED isolation level

# Full Snapshot Isolation


---

```
ALTER DATABASE MyDatabase  
SET ALLOW_SNAPSHOT_ISOLATION ON
```

- operations see the most recent committed data as of the beginning of their transaction =>
  - snapshot of the data at the transaction level
  - consistent reads at the transaction level
  - SNAPSHOT isolation level

# Row-Level Versioning

- each version of a row is marked with the *transaction sequence number (XSN)* of the transaction that changed the row
- all versions are kept in a linked list:



# Row-Level Versioning

---

- advantages:
  - increased concurrency level
  - positive impact - triggers / indexes
- drawbacks:
  - monitoring the usage of the *tempdb* database => extra management requirements
  - update operations – slower
  - read operations – also affected (browsing the linked lists)



# Row-Level Versioning

---

- drawbacks
  - solves the writer-reader conflict, but simultaneous writers are still not allowed

# Triggers & RLV

---

- triggers can access 2 *pseudo-tables*:
  - the *deleted* table – contains removed rows or previous versions of updated rows
  - the *inserted* table – contains added rows or new versions of updated rows
- row versioning is used to create the *inserted* / *deleted* tables
- before SQL Server 2005:
  - the *deleted* table was created using the transaction log – affected performance

# Index Creation & RLV

---

- index creation / rebuilding in previous versions of SQL Server:
  - *clustered index* - table exclusively locked, data entirely inaccessible
  - *non-clustered index* - index not available
- using RLV:
  - indexes are created and rebuilt online
  - all requests are processed on versioned data

# Isolation Levels and Concurrency Anomalies

concurrency problem / isolation level	read uncommitted	read committed locking	read committed snapshot	repeatable read	snapshot	serializable
dirty reads	yes	no	no	no	no	no
unrepeatable reads	yes	yes	yes	no	no	no
phantoms	yes	yes	yes	yes	no	no
update conflicts	no	no	no	no	yes	no
concurrency model	pessimistic	pessimistic	optimistic	pessimistic	optimistic	pessimistic

# Query Governor and DBCC LOG

- SET QUERY\_GOVERNOR\_COST\_LIMIT *value*
  - *value*
    - the longest time in which a query can run
  - queries with an estimated cost greater than *value* are not allowed to run
  - value 0 – all queries are allowed to run
- DBCC LOG – transaction log
  - DBCC LOG (<DBname>,<Output>)
  - Output - level of detail (0-4)

# PIVOT / UNPIVOT

- change a table-valued expression into another table
- PIVOT rotates a table-valued expression; it transforms the unique values in one column in the expression into multiple columns in the output; aggregations are performed where necessary on any remaining column values that are required in the output
- UNPIVOT performs the opposite operation; it rotates columns in a table-valued expression into column values

# PIVOT

```
SELECT <non-pivoted column>,  
    [first pivoted column] AS <column name>,  
    [second pivoted column] AS <column name>,  
    ...  
    [last pivoted column] AS <column name>  
FROM  
    (<SELECT query that produces the data>) AS  
        <source query alias>  
PIVOT  
(  
    <aggregation function>(<column being aggregated>)  
FOR  
    [<column that contains values that become column headers>]  
    IN ( [first pivoted column], [second pivoted column],  
        ... [last pivoted column])  
) AS <alias for the pivot table>  
<optional ORDER BY clause>;
```

# Recap - The OUTPUT Clause

- provides access to *inserted*, *updated*, *deleted* records
- can implement certain functionalities which can otherwise be performed only via triggers

```
UPDATE Courses
SET cname = 'Database Management Systems'
OUTPUT inserted.cid, deleted.cname, inserted.cname,
      GETDATE(), SUSER_SNAME()
INTO CourseChanges
WHERE cid = 'DB2'
```



# Recap - The MERGE Statement

---

■ a source table is compared with a target table; INSERT, UPDATE, DELETE statements can be executed based on the result of the comparison, i.e., INSERT / UPDATE / DELETE operations can be executed on the target table based on the result of a join with the source table

# Recap - MERGE – General Syntax

---

```
MERGE TargetTable AS Target
USING SourceTable AS Source
ON (Search terms)
WHEN MATCHED THEN
    UPDATE SET
    or
    DELETE
WHEN NOT MATCHED [BY TARGET] THEN
    INSERT
WHEN NOT MATCHED BY SOURCE THEN
    UPDATE SET
    or
    DELETE
```

# Recap - MERGE example

## Books table

	BookID	Title	Author	ISBN	Pages
1	1	In Search of Lost Time	Marcel Proust	NULL	NULL
2	2	In Search of Lost Time	NULL	NULL	350
3	3	In Search of Lost Time	NULL	9789731246420	NULL

# Recap - MERGE example

---

MERGE Books  
USING

```
(SELECT MAX(BookID) BookID, Title, MAX(Author)  
Author, MAX(ISBN) ISBN, MAX(Pages) Pages  
FROM Books  
GROUP BY Title
```

```
) MergeData ON Books.BookID = MergeData.BookID  
WHEN MATCHED THEN  
    UPDATE SET Books.Title = MergeData.Title,  
    Books.Author = MergeData.Author,  
    Books.ISBN = MergeData.ISBN,  
    Books.Pages = MergeData.Pages  
WHEN NOT MATCHED BY SOURCE THEN DELETE;
```

# Recap - MERGE example

---

	BookID	Title	Author	ISBN	Pages
1	3	In Search of Lost Time	Marcel Proust	9789731246420	350

# Seminar 5

---

## Performance Tuning in SQL Server

# Query Tuning Methodology

---

- identify waits (bottleneck) at the server level
  - locks
  - transaction log
  - I/O
  - etc.
- correlate waits with queues
- drill down to database / file level
- drill down to process level
- tune problematic queries
- \* DMVs - dynamic management views

# Identify Waits

---

## ■ `sys.dm_os_wait_stats`

### ■ returned table

- `wait_type`
  - resource waits (locks, latches, network, I/O), queue waits, external waits
- `waiting_tasks_count`
- `wait_time_ms`
- `max_wait_time_ms`
- `signal_wait_time_ms`

## ■ reset DMV values

- `DBCC SQLPERF ('sys.dm_os_wait_stats', CLEAR);`



# Correlate Waits with Queues

---

- **sys.dm\_os\_performance\_counters**
  - *object\_name* - the category of the counter
  - *counter\_name* - the name of the counter
  - *instance\_name* - the name of the specific instance of the counter; often contains the name of the database
  - *cntr\_value* - the current value of the counter
  - *cntr\_type* - the type of the counter (as defined by the Windows performance architecture)

# Correlate Waits with Queues

---

- **sys.dm\_os\_performance\_counters**
- > 500 counters: Access Methods, User Settable, Buffer Manager, Broker Statistics, SQL Errors, Latches, Buffer Partition, SQL Statistics, Locks, Buffer Node, Plan Cache, Cursor Manager by Type, Memory Manager, General Statistics, Databases, Catalog Metadata, Broker Activation, Broker/DBM Transport, Transactions, Cursor Manager Total, Exec Statistics, Wait Statistics, etc.
- *cntr\_type* = 65792 → *cntr\_value* contains the actual value

# Correlate Waits with Queues

---

- `sys.dm_os_performance_counters`
- *cntr\_type* = 537003264 → *cntr\_value* contains real-time results, which are divided by a “base” to obtain the actual value; by themselves, they are useless
  - to get a ratio: divide by a “base” value
  - to get a percentage: multiply the result by 100.0

# Correlate Waits with Queues

---

- `sys.dm_os_performance_counters`
- `cntr_type = 272696576`
  - time-based, cumulative counters
  - a secondary table can be used to log intermediate values

# Correlate Waits with Queues

- `sys.dm_os_performance_counters`
- `cntr_type = 1073874176` and `cntr_type = 1073939712` → poll both the value (1073874176) and the base value (1073939712)
- poll both values again (e.g., after 15 seconds) ☺
- to obtain the desired result, compute:
$$\text{UnitsPerSec} = (\text{cv2} - \text{cv1}) / (\text{bv2} - \text{bv1}) / 15.0$$

# Drill Down to Database / File Level

---

## ■ `sys.dm_io_virtual_file_stats`

- returns I/O information about *data files* and *log files*

## ■ parameters

### ■ `database_ID`

- NULL = all databases
- useful function: DB\_ID

### ■ `file_ID`

- NULL = all files
- useful function: FILE\_IDEX

# Drill Down to Database / File Level

---

## ■ `sys.dm_io_virtual_file_stats`

### ■ returned table

- `database_ID`
- `file_ID`
- `sample_ms` - # of milliseconds since the computer was started
- `num_of_reads` - number of reads issued on the file
- `num_of_bytes_read` - number of bytes read on the file
- `io_stall_read_ms` - total time users waited for reads issued on the file

# Drill Down to Database / File Level

---

## ■ `sys.dm_io_virtual_file_stats`

### ■ returned table

- `num_of_writes` - number of writes
- `num_of_bytes_written` - total number of bytes written to the file
- `io_stall_write_ms` - total time users waited for writes to be completed on the file
- `io_stall` - total time users waited for the completion of I/O operations (ms)
- `file_handle`



# Drill Down to the Process Level

---

- a filter on duration / I/O only isolates individual processes (batch / proc / query)
- aggregate performance information by query pattern
  - patterns can be easily identified when using stored procedures
  - when one doesn't use stored procedures:
    - quick and dirty approach: LEFT(query string, n)
    - use a parser to identify the query pattern

# Indexes

---

- one of the major factors influencing query performance
  - impact on: filtering, joins, sorting, grouping; blocking and deadlock avoidance, etc.
  - effect on modifications: positive effect (locating the rows); negative effect (cost of modifying the index)
- understanding indexes and their internal mechanisms
  - clustered/nonclustered, single/multicolumn, indexed views, indexes on computed columns, covering scenarios, intersection

# Indexes

---

- one should carefully judge whether additional index maintenance costs are justified by improvements in query performance
  - take into account the environment and the ratio between SELECT queries and data modifications
- *multicolumn indexes*
  - tend to be more useful than single-column indexes
  - the query optimizer is more likely to use such indexes to cover a query

# Indexes

---

- *indexed views* come with a higher maintenance cost than standard indexes
  - mandatory option
    - WITH SCHEMABINDING

# Tools to Analyze Query Performance

---

- graphical execution plan
- STATISTICS IO - scan count, logical reads, physical reads, *read-ahead* reads
- STATISTICS TIME - duration and net CPU time
- SHOWPLAN\_TEXT - SQL Server returns detailed information about how the statements are executed
- SHOWPLAN\_ALL - SQL Server returns detailed information about how the statements are executed, provides estimates of the resource requirements

# Tools to Analyze Query Performance

---

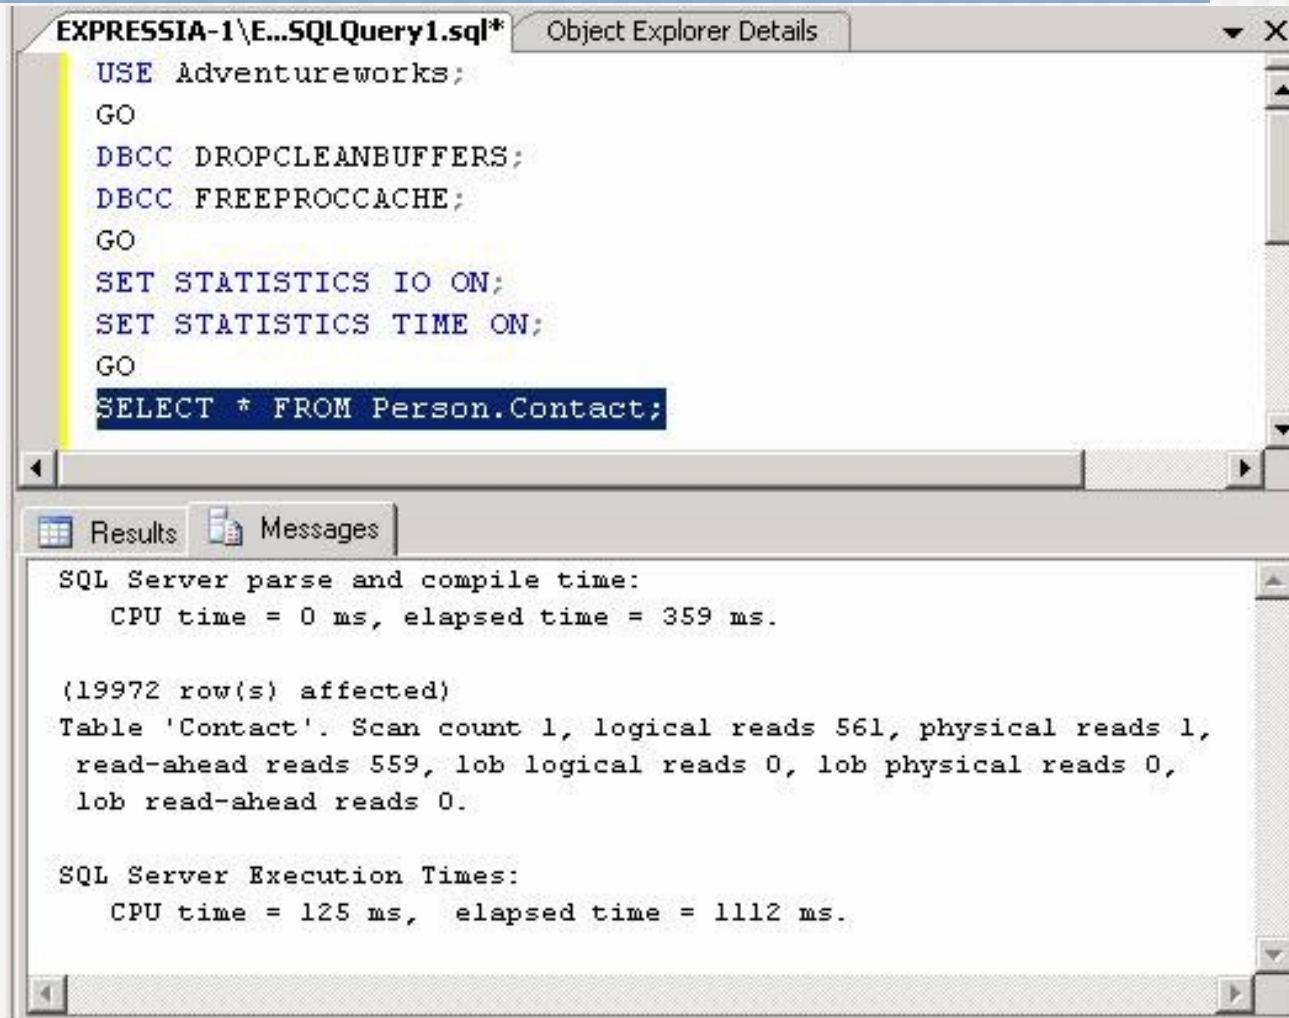
- STATISTICS PROFILE - a profile of the query execution
- STATISTICS XML - actual plan information in XML format
- SHOWPLAN\_XML - estimated plan information in XML format

# Query Optimization

---

- evaluating execution plans
  - sequences of physical/logical operations
- optimization - factors
  - search predicate
  - tables involved in joins
  - join conditions
  - result set size
  - list of indexes
- goal - avoid worst query plans
- SQL Server uses a *cost-based* query optimizer

# STATISTICS IO and STATISTICS TIME



The screenshot shows a SQL Server Enterprise Manager window titled 'EXPRESSIA-1\E...SQLQuery1.sql\*'. The query window contains the following T-SQL code:

```
USE Adventureworks;  
GO  
DBCC DROPCLEANBUFFERS;  
DBCC FREEPROCCACHE;  
GO  
SET STATISTICS IO ON;  
SET STATISTICS TIME ON;  
GO  
SELECT * FROM Person.Contact;
```

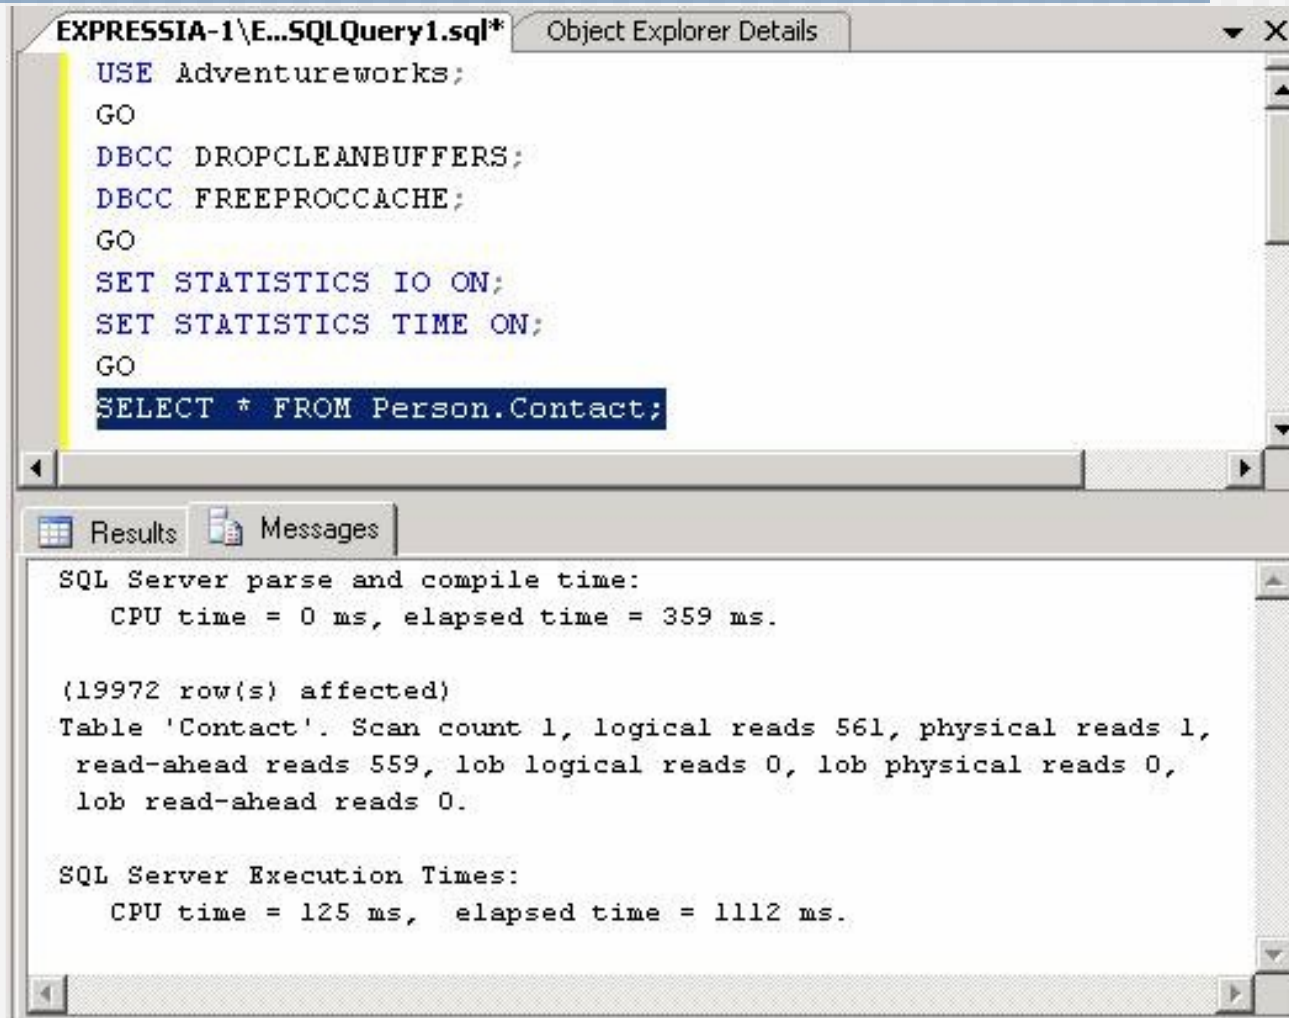
The 'Results' tab is selected, displaying the following output:

```
SQL Server parse and compile time:  
    CPU time = 0 ms, elapsed time = 359 ms.  
  
(19972 row(s) affected)  
Table 'Contact'. Scan count 1, logical reads 561, physical reads 1,  
    read-ahead reads 559, lob logical reads 0, lob physical reads 0,  
    lob read-ahead reads 0.  
  
SQL Server Execution Times:  
    CPU time = 125 ms,  elapsed time = 1112 ms.
```

- DBCC DROPCLEANBUFFERS – clear data from the cache
- DBCC FREEPROCCACHE – clear execution plans from the cache



# STATISTICS IO and STATISTICS TIME



The screenshot shows a SQL Server Enterprise Manager window titled 'EXPRESSIA-1\E...SQLQuery1.sql\*'. The query window contains the following SQL code:

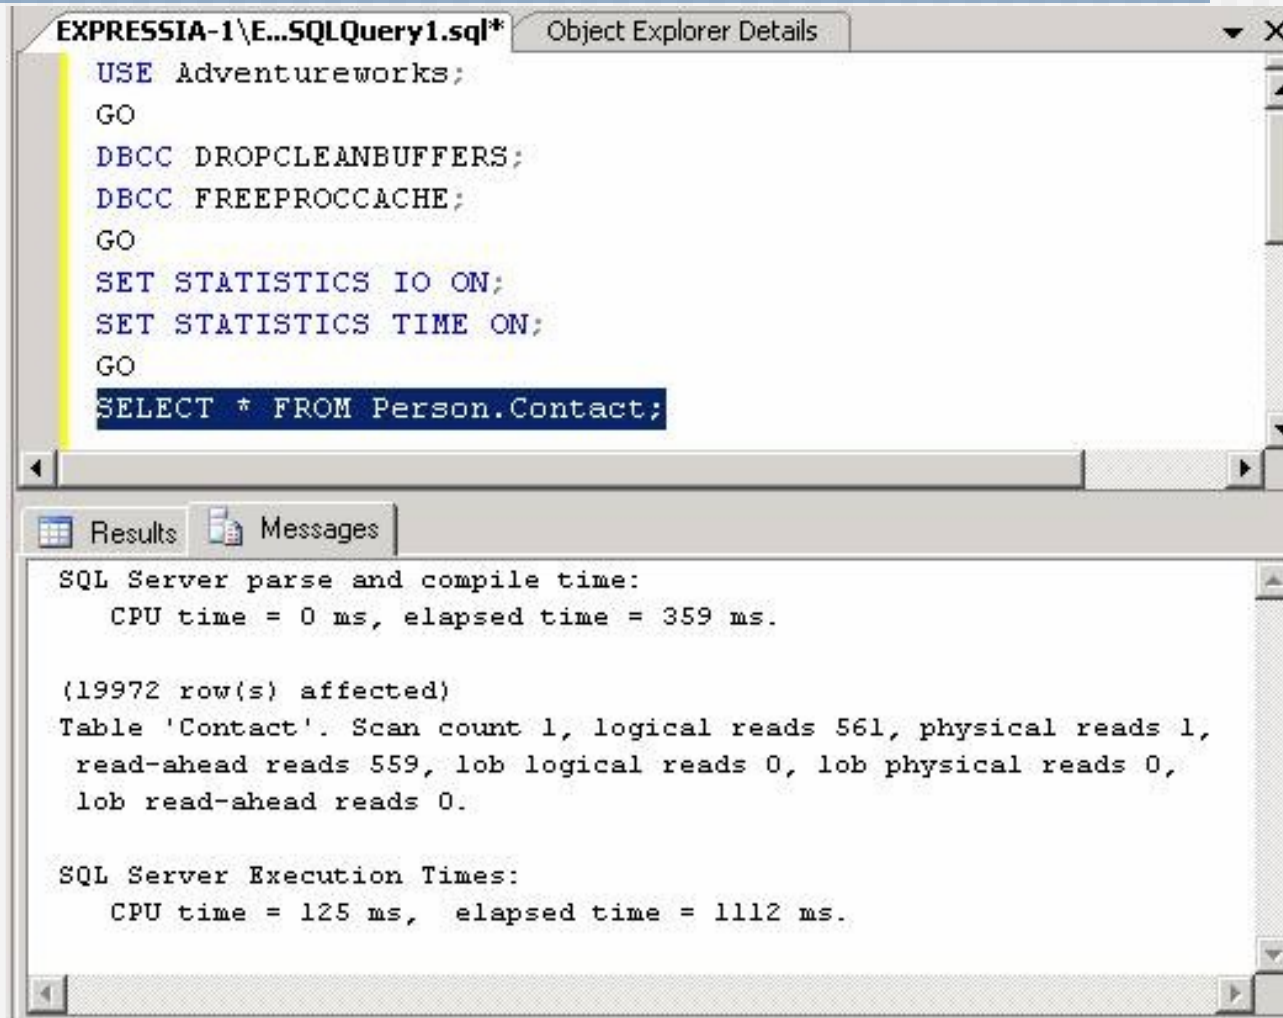
```
USE Adventureworks;  
GO  
DBCC DROPCLEANBUFFERS;  
DBCC FREEPROCCACHE;  
GO  
SET STATISTICS IO ON;  
SET STATISTICS TIME ON;  
GO  
SELECT * FROM Person.Contact;
```

The 'Results' tab is selected, displaying the following output:

```
SQL Server parse and compile time:  
    CPU time = 0 ms, elapsed time = 359 ms.  
  
(19972 row(s) affected)  
Table 'Contact'. Scan count 1, logical reads 561, physical reads 1,  
    read-ahead reads 559, lob logical reads 0, lob physical reads 0,  
    lob read-ahead reads 0.  
  
SQL Server Execution Times:  
    CPU time = 125 ms,  elapsed time = 1112 ms.
```

- *CPU time, elapsed time* – CPU time and elapsed time to parse and compile the query, and to execute it

# STATISTICS IO and STATISTICS TIME



The screenshot shows a SQL Server Enterprise Manager window titled 'EXPRESSIA-1\E...SQLQuery1.sql\*'. The query window contains the following SQL code:

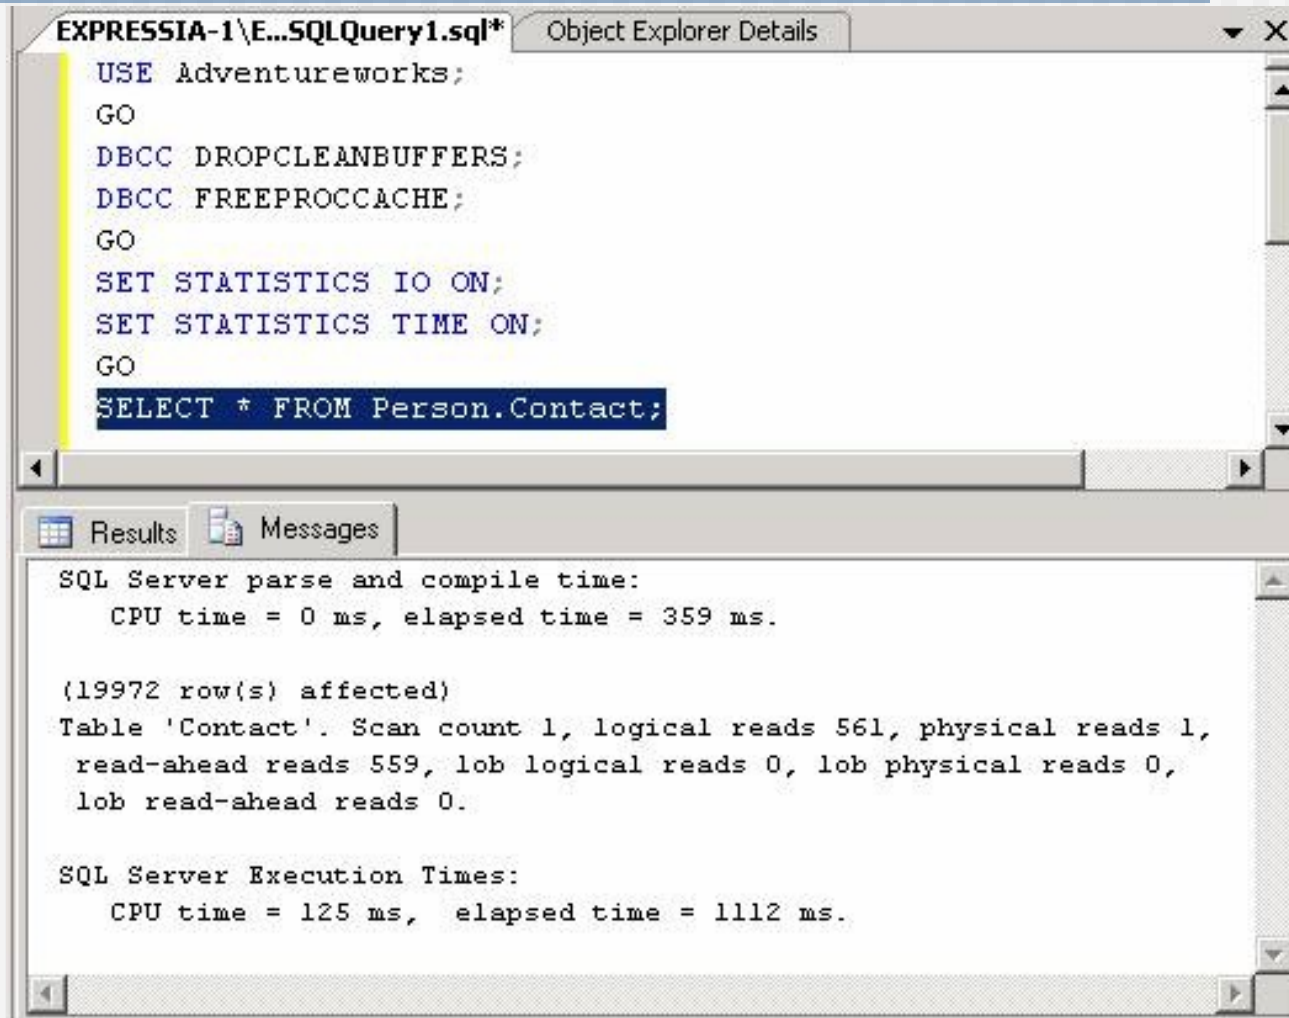
```
USE Adventureworks;  
GO  
DBCC DROPCLEANBUFFERS;  
DBCC FREEPROCCACHE;  
GO  
SET STATISTICS IO ON;  
SET STATISTICS TIME ON;  
GO  
SELECT * FROM Person.Contact;
```

The 'Results' tab is selected, displaying the following output:

```
SQL Server parse and compile time:  
    CPU time = 0 ms, elapsed time = 359 ms.  
  
(19972 row(s) affected)  
Table 'Contact'. Scan count 1, logical reads 561, physical reads 1,  
    read-ahead reads 559, lob logical reads 0, lob physical reads 0,  
    lob read-ahead reads 0.  
  
SQL Server Execution Times:  
    CPU time = 125 ms,  elapsed time = 1112 ms.
```

- *physical reads* - number of pages read from the disk
- *read-ahead reads* - number of pages placed in the cache for the query

# STATISTICS IO and STATISTICS TIME



The screenshot shows a SQL Server Enterprise Manager window titled 'EXPRESSIA-1\E...SQLQuery1.sql\*'. The query window contains the following T-SQL code:

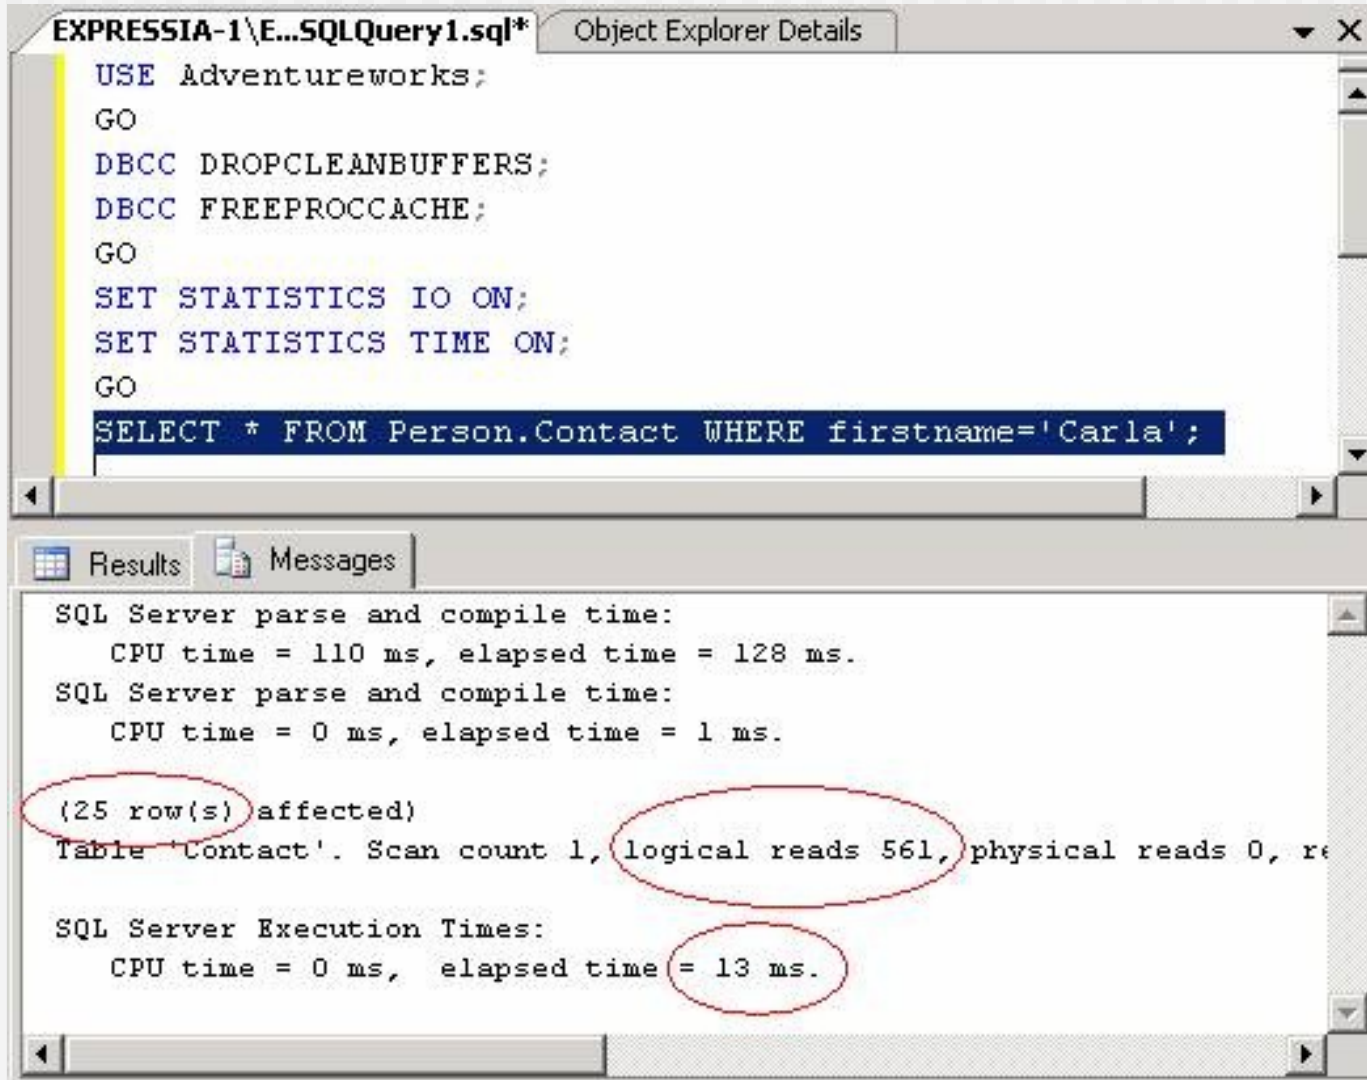
```
USE Adventureworks;  
GO  
DBCC DROPCLEANBUFFERS;  
DBCC FREEPROCCACHE;  
GO  
SET STATISTICS IO ON;  
SET STATISTICS TIME ON;  
GO  
SELECT * FROM Person.Contact;
```

The 'Results' tab is selected, displaying the following output:

```
SQL Server parse and compile time:  
    CPU time = 0 ms, elapsed time = 359 ms.  
  
(19972 row(s) affected)  
Table 'Contact'. Scan count 1, logical reads 561, physical reads 1,  
    read-ahead reads 559, lob logical reads 0, lob physical reads 0,  
    lob read-ahead reads 0.  
  
SQL Server Execution Times:  
    CPU time = 125 ms,  elapsed time = 1112 ms.
```

- *scan count* - number of seeks or scans after reaching the leaves
- *logical reads* - number of pages read from the data cache

# STATISTICS IO and STATISTICS TIME



The screenshot shows a SQL Server Enterprise Manager window with a query window titled 'EXPRESSIA-1\E...SQLQuery1.sql\*' and an 'Object Explorer Details' pane. The query window contains the following SQL code:

```
USE Adventureworks;  
GO  
DBCC DROPCLEANBUFFERS;  
DBCC FREEPROCCACHE;  
GO  
SET STATISTICS IO ON;  
SET STATISTICS TIME ON;  
GO  
SELECT * FROM Person.Contact WHERE firstname='Carla';
```

Below the query window, the 'Results' pane shows the execution statistics for the query. The statistics are as follows:

- SQL Server parse and compile time:  
CPU time = 110 ms, elapsed time = 128 ms.
- SQL Server parse and compile time:  
CPU time = 0 ms, elapsed time = 1 ms.
- (25 row(s) affected)
- Table 'Contact'. Scan count 1, logical reads 561, physical reads 0, re
- SQL Server Execution Times:  
CPU time = 0 ms, elapsed time = 13 ms.

Red circles highlight the following values in the results:

- (25 row(s) affected)
- logical reads 561
- elapsed time = 13 ms.

# STATISTICS IO and STATISTICS TIME

---

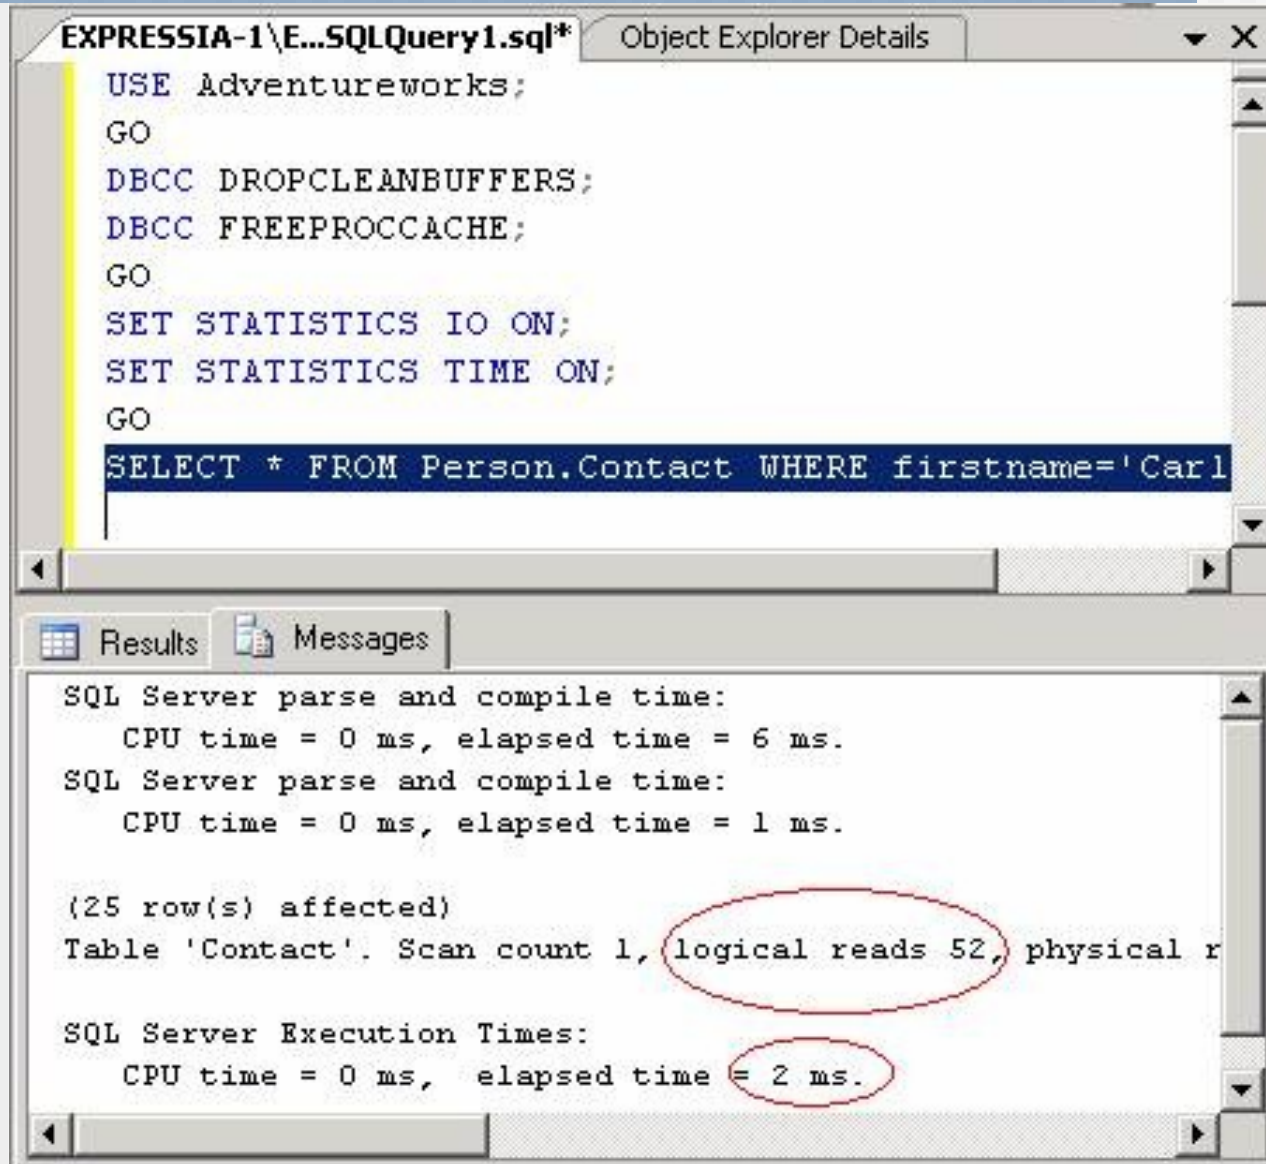
```
USE AdventureWorks
```

```
GO
```

```
CREATE NONCLUSTERED INDEX IDX_FirstName  
    ON Person.Contact (FirstName ASC)
```

```
GO
```

# STATISTICS IO and STATISTICS TIME



The screenshot shows a SQL Server Enterprise Manager window with a query editor and a results pane. The query editor contains the following SQL code:

```
USE Adventureworks;
GO
DBCC DROPCLEANBUFFERS;
DBCC FREEPROCCACHE;
GO
SET STATISTICS IO ON;
SET STATISTICS TIME ON;
GO
SELECT * FROM Person.Contact WHERE firstname='Carl'
```

The results pane shows the execution statistics for the query:

```
SQL Server parse and compile time:
    CPU time = 0 ms, elapsed time = 6 ms.
SQL Server parse and compile time:
    CPU time = 0 ms, elapsed time = 1 ms.

(25 row(s) affected)
Table 'Contact'. Scan count 1, logical reads 52, physical r

SQL Server Execution Times:
    CPU time = 0 ms, elapsed time = 2 ms.
```

The values "logical reads 52" and "elapsed time = 2 ms." are circled in red in the original image.




# Graphical Execution Plan

```
USE AdventureWorks
GO
SELECT COUNT(*) cRows
FROM HumanResources.Shift;
GO
```

Results Execution plan

Query 1: Query cost (relative to the batch): 100%

SELECT COUNT(\*) cRows FROM HumanResources.Shift;



# SHOWPLAN\_ALL

```
SET SHOWPLAN_ALL ON;  
GO  
SELECT COUNT(*) cRows  
FROM HumanResources.Shift;  
GO  
SET SHOWPLAN_ALL OFF;  
GO
```



Results

StmtText

```
-----  
SELECT COUNT(*) cRows  
FROM HumanResources.Shift;  
|--Compute Scalar(DEFINE:([Expr1003]=CONVERT_IMPLICIT(int,[Expr1004],0)))  
--Stream Aggregate(DEFINE:([Expr1004]=Count(*)))  
--Index Scan(OBJECT:([master].[HumanResources].[Shift].[AK_Shift])  
  
(4 row(s) affected)
```

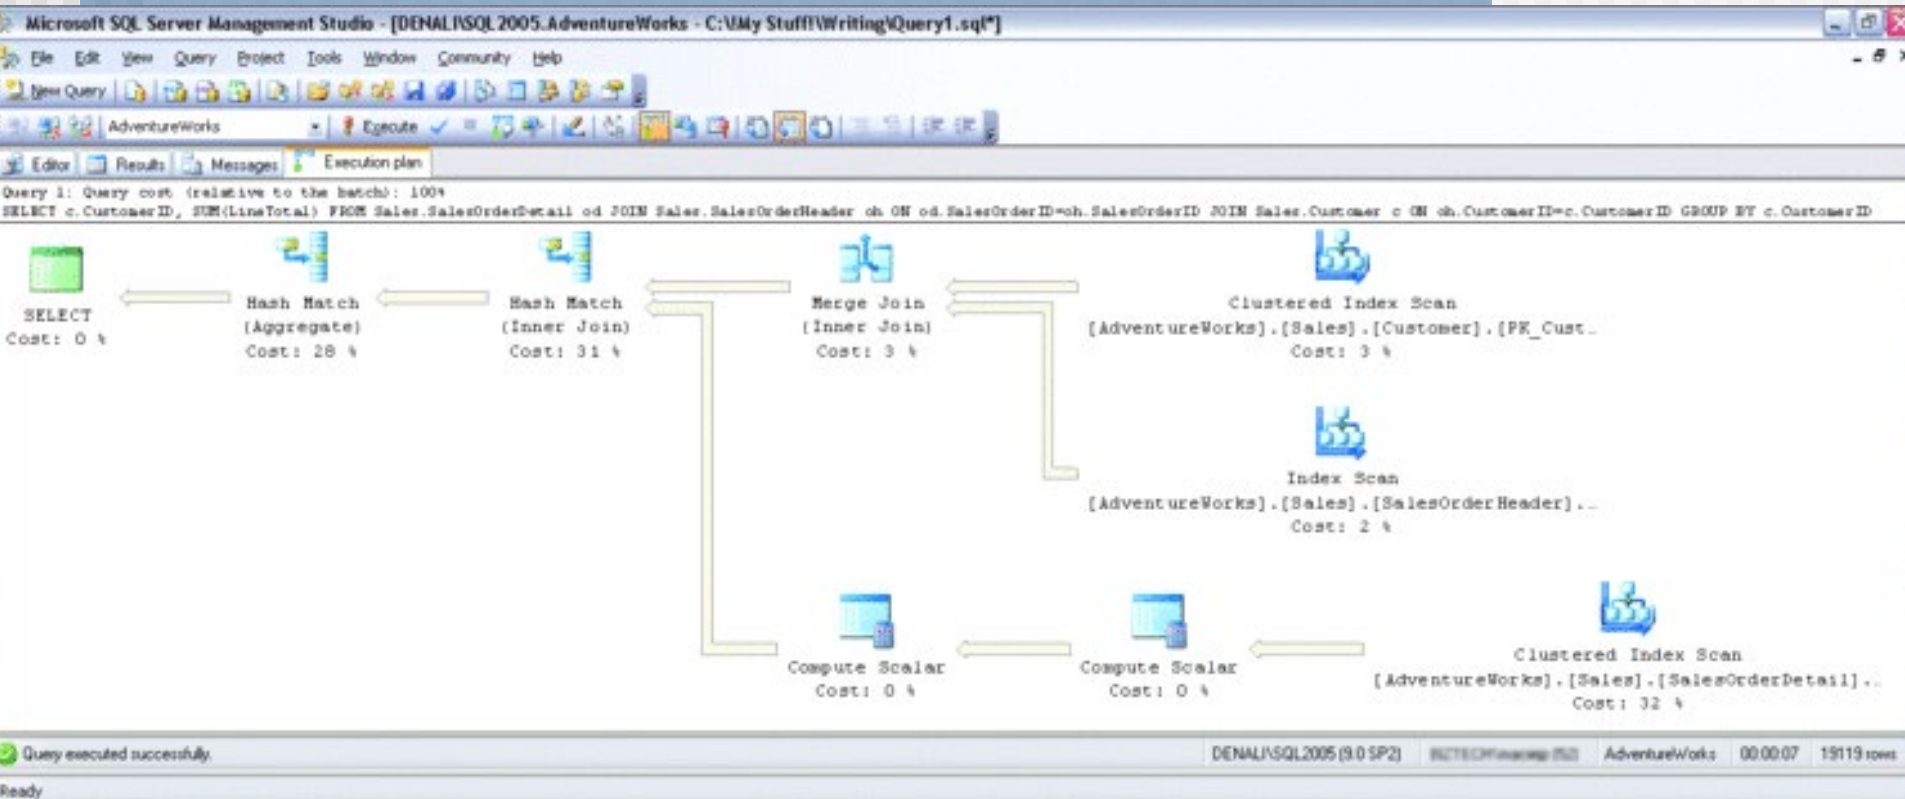


# Graphical Execution Plan

---

```
SELECT c.CustomerID, SUM(LineTotal)
FROM Sales.SalesOrderDetail od
      JOIN Sales.SalesOrderHeader oh ON
          od.SalesOrderID = oh.SalesOrderID
      JOIN Sales.Customer c ON
          oh.CustomerID = c.CustomerID
GROUP BY c.CustomerID
```

# Graphical Execution Plan



# Graphical Execution Plan



SELECT

Cost: 0 %



Hash Match  
(Aggregate)

SELECT

<b>Cached plan size</b>	40 B
<b>Degree of Parallelism</b>	0
<b>Memory Grant</b>	812
<b>Estimated Operator Cost</b>	0 (0%)
<b>Estimated Subtree Cost</b>	3,31365
<b>Estimated Number of Rows</b>	19045

## Statement

```
SELECT c.CustomerID, SUM(LineTotal)
FROM Sales.SalesOrderDetail od JOIN
Sales.SalesOrderHeader oh
ON od.SalesOrderID=oh.SalesOrderID
JOIN Sales.Customer c ON
oh.CustomerID=c.CustomerID
GROUP BY c.CustomerID
```



Clustered Index Scan

[AdventureWorks].[Sales].[SalesOrderDetail]...

Cost: 32 %

## Clustered Index Scan

Scanning a clustered index, entirely or only a range.

<b>Physical Operation</b>	Clustered Index Scan
<b>Logical Operation</b>	Clustered Index Scan
<b>Actual Number of Rows</b>	121317
<b>Estimated I/O Cost</b>	0,915718
<b>Estimated CPU Cost</b>	0,133606
<b>Estimated Operator Cost</b>	1,04932 (32%)
<b>Estimated Subtree Cost</b>	1,04932
<b>Estimated Number of Rows</b>	121317
<b>Estimated Row Size</b>	29 B
<b>Actual Rebinds</b>	0
<b>Actual Rewinds</b>	0
<b>Ordered</b>	False
<b>Node ID</b>	8

## Object

[AdventureWorks].[Sales].[SalesOrderDetail].  
[PK\_SalesOrderDetail\_SalesOrderID\_SalesOrderDetailID]  
[od]

## Output List

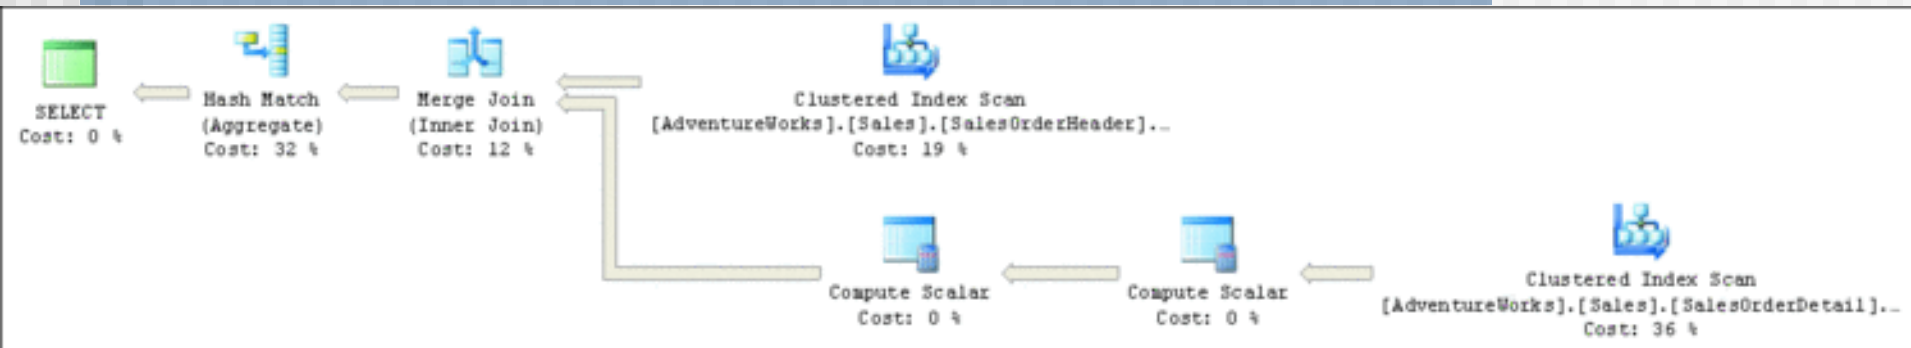
[AdventureWorks].[Sales].  
[SalesOrderDetail].SalesOrderID; [AdventureWorks].  
[Sales].[SalesOrderDetail].OrderQty; [AdventureWorks].  
[Sales].[SalesOrderDetail].UnitPrice; [AdventureWorks].  
[Sales].[SalesOrderDetail].UnitPriceDiscount

# Graphical Execution Plan

---

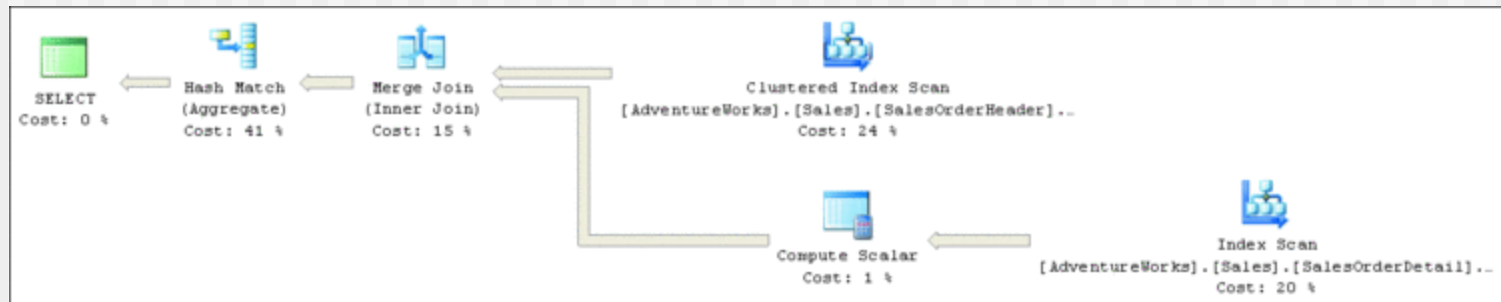
```
SELECT oh.CustomerID, SUM(LineTotal)
FROM Sales.SalesOrderDetail od
      JOIN Sales.SalesOrderHeader oh ON
          od.SalesOrderID=oh.SalesOrderID
GROUP BY oh.CustomerID
```

# Graphical Execution Plan



# Graphical Execution Plan

```
CREATE INDEX IDX_OrderDetail_OrderID_TotalLine  
ON Sales.SalesOrderDetail (SalesOrderID)  
INCLUDE (LineTotal)
```



Model for the Practical Exam

I. *Customers* is a table in a SQL Server database with schema *Customers*[*CustomerID*, *FirstName*, *LastName*, *City*, *DateOfBirth*]. The primary key is underlined.

*CustomerID* is the search key of the clustered index on *Customers*. The table doesn't have any other indexes.

Consider the interleaved execution below. There are no other concurrent transactions. The value of *City* for the customer with *CustomerID* 2 is *Timisoara* when T1 begins execution.

Answer questions 1-3 (each question has at least one correct answer).

T1	T2
BEGIN TRAN SELECT City FROM Customers WHERE CustomerID = 2	
	BEGIN TRAN UPDATE Customers SET City = 'Cluj-Napoca' WHERE CustomerID = 2
UPDATE Customers SET City = 'Bucuresti' WHERE CustomerID = 2	
	ROLLBACK TRAN
COMMIT TRAN	

time

1. T1 and T2 run under READ UNCOMMITTED. After the *COMMIT TRAN* statement in T1, the *City* value for the customer with *CustomerID* 2 is:

- Timisoara*
- Cluj-Napoca*
- Bucuresti*
- NULL
- None of the above answers is correct.

2. T1 runs under READ COMMITTED and T2 under REPEATABLE READ. After the *COMMIT TRAN* statement in T1, the *City* value for the customer with *CustomerID* 2 is:

- Timisoara*
- Cluj-Napoca*
- Bucuresti*
- NULL
- None of the above answers is correct.

3. T1 runs under REPEATABLE READ and T2 runs under READ COMMITTED. Then:

- T1 doesn't acquire a shared lock for its SELECT statement.
- T1 acquires a shared lock for its SELECT statement.
- T2 needs an exclusive lock for its UPDATE statement.
- T1 needs an exclusive lock for its UPDATE statement.
- None of the above answers is correct.

**II.** Create a database for a MiniFacebook system. The entities of interest to the problem domain are: *Users*, *Pages*, *Likes*, *Categories*, *Posts*, and *Comments*. Each user has a name, current city and date of birth. A user can like multiple pages. The system stores the date of each like. A page has a name and a category, e.g., *sports*, *movies*, *music*, etc. A category also has a category description. Users write posts and comment on existing posts. A user's post has a date, text, and number of shares. A comment is anonymous, has a text, a date, and a flag indicating whether it's a top comment for the corresponding post.

1. Write an SQL script that creates the corresponding relational data model.

2. Create a Master/Detail Form that allows one to display the posts for a given user, to carry out <insert, update, delete> operations on the posts of a given user. The form should have a *DataGridView* named *dgvUsers* to display the users, a *DataGridView* named *dgvPosts* to display all the posts of the selected user, and a button for saving added / deleted / modified posts. You must use the following classes: *DataSet*, *SqlDataAdapter*, *BindingSource*.

3. Create a scenario that reproduces the non-repeatable read concurrency issue on this database. Explain why the non-repeatable read occurs, and describe a solution to prevent this concurrency issue. Don't use stored procedures.

I. 1	1p
2	1p
3	1p
II. 1	2p
2	2p
3	2p
	1p of